



Introduction to Generative AI

2. Prompt Engineering




Text Processing

- This includes code
- Machine learning algorithms use linear algebra, statistical methods, and other mathematical techniques to train models
 - Non-numeric training data has to be converted into a numerical form
 - Language data is a string of characters
 - Tokenization converts the string of characters into a string of words
 - Vectorization converts tokens into numerical vectors called parameters




Tokenization

- Tokenization is the process of splitting text into smaller units called tokens
- Tokens can be
 - Words
 - Subwords
 - Characters
 - Special symbols
 - Example
 - Text : “Transformers are powerful.”
 - Tokenized: [“Transformers”, “are”, “powerful”, “.”]




Vector Embedding

- Because LLMs work on numbers, the tokens need to be converted into vectors
- Color terms example
 - Can be turned into vectors with three dimensions (red, green, blue) with values from 0 to 255
 - We can then determine similarity of colors by using vector distances
 - *Orange (255,165,0) is closer to red (255,0,0) than blue (0,0,255)*
 - There are different ways of vectorizing colors, most use three dimensions
 - *Instead of red/green/blue, we could use magenta/yellow/cyan*




Vector Embedding

- We could do something similar for other words by choosing a set of meaningful dimensions or parameters
- The selection of parameters is usually pre-selected in classical machine learning
 - Usually those that should be most useful in building the model
 - For example, income is more useful for predicting credit worthiness and shoe size is less useful




Choosing Features

- Problem, choosing the right features
- Solution: Use a neural net
 - It will do automatic feature extraction based on patterns in the data
 - But the choice of features is opaque to us
 - The more features we use, the better the performance
 - Most LLMs use billions or trillions of features




Training LLMs

- Pretraining
 - The model learns general language understanding from a vast corpus of text data
 - This often uses unsupervised learning techniques
 - Produces a base of foundation model
- Fine-tuning
 - Model is further trained on specific tasks or domains, adapting its knowledge to particular applications
 - For example, Claude code and copilot are trained on code




Training LLMs

- Pre-training is usually done on
 - Massive internet-scale text corpora
 - Books, articles, code, websites, etc.
 - Trillions of tokens
- A base model is is not “chat ready”
 - Prompt: “Explain quantum mechanics”
 - A base model might
 - *Continue like a textbook*
 - *Change tone mid-way*
 - *Provide incomplete structure*
 - *Produce unsafe or unfiltered content*
 - It has no alignment, which is
 - *Shaping an AI system's behavior so that its outputs are consistent with human values, intentions, and safety expectations*



Supervised Fine-Tuning

- SFT trains the model on curated prompt–response pairs
- The model learns to
 - Respond directly to instructions
 - Format answers cleanly
 - Follow conversational patterns
- After SFT, the model
 - Follows instructions more reliably
 - Adopts assistant-style tone
 - Structures answers clearly



Emergent Behavior

- Emergent behavior means
 - A capability appears suddenly once the model reaches a sufficiently large scale
 - Examples observed in large models
 - *In-context learning (learning from examples in the prompt)*
 - *Multi-step reasoning*
 - *Chain-of-thought reasoning*
 - *Arithmetic competence*
 - *Complex code generation*
 - *Tool-use planning behavior*
- Smaller models may fail completely at these tasks
 - Larger models suddenly perform well

Predictability and Surprise in Large Generative Models

FACCT '22, June 21–24, 2022, Seoul, Republic of Korea




Fig. 2 Three examples of abrupt specific capability scaling described in Section 2.2, based on three different models: GPT-3 (blue), Gopher (orange), and a Google language model (green). **(Left)** 3-Digit addition with GPT-3 [11]. **(Middle)** Language understanding with GPT-3 and Gopher [62]. **(Right)** Program synthesis with Google language models [4].



Emergent Behavior

- Emergent behavior are threshold effects
- Certain reasoning abilities require
 - Enough parameters
 - Enough training data
 - Enough depth
- Below a threshold, this capability is weak and unusable
- Above a threshold, this capability becomes usable

Predictability and Surprise in Large Generative Models

FAccT '22, June 21–24, 2022, Seoul, Republic of Korea




Fig. 2 Three examples of abrupt specific capability scaling described in Section 2.2, based on three different models: GPT-3 (blue), Gopher (orange), and a Google language model (green). **(Left)** 3-Digit addition with GPT-3 [11]. **(Middle)** Language understanding with GPT-3 and Gopher [62]. **(Right)** Program synthesis with Google language models [4].



Emergent Behavior

- As models grow
 - Representations become richer
 - Abstractions become more layered
 - Internal reasoning chains become deeper
- Deep stacking and larger scale allows
 - Multi-step transformations
 - More complex internal “thought-like” structures

Predictability and Surprise in Large Generative Models

FAccT '22, June 21–24, 2022, Seoul, Republic of Korea





Fig. 2 Three examples of abrupt specific capability scaling described in Section 2.2, based on three different models: GPT-3 (blue), Gopher (orange), and a Google language model (green). **(Left)** 3-Digit addition with GPT-3 [11]. **(Middle)** Language understanding with GPT-3 and Gopher [62]. **(Right)** Program synthesis with Google language models [4].




The Core Problem

- When a model generates output
 - It samples from a smoothed probability distribution
 - Rare events become even rarer
 - Common patterns become more dominant
- If a new model is trained on generated output
 - It learns the distorted smoothed distribution
 - Rare but important patterns disappear
 - Diversity shrinks
- After multiple generations
 - The distribution converges to a very small set of outputs
 - This is called model collapse




The Core Problem

- The internet is increasingly filled with
 - AI-generated articles
 - AI-generated code
 - AI-generated forum posts
 - AI-generated marketing content
- Future models may scrape
 - AI generated content that was generated by earlier models
 - Training data becomes contaminated with LLM generated content
 - Similar in concept to genetic anomalies introduce by inbreeding in a population




Hallucinations

- A hallucination occurs when a model generates information that is false, fabricated, or unsupported
 - Examples
 - *Inventing academic citations*
 - *Creating fake legal cases*
 - *Fabricating statistics*
 - *Misstating historical facts*
 - *Generating imaginary test cases*
 - *Writing code that compiles but does not execute correctly*
- The model does not “know” it is hallucinating




Hallucinations

- LLMs are trained to
 - Predict likely next tokens
 - Not verify truth
- They model probability, not reality
 - If a plausible-sounding answer fits the language or code pattern
 - The model may produce it
 - Even if it is incorrect



Hallucinations

- Why this is dangerous
 - High confidence can mislead users
 - Hard to detect without external validation
 - Especially risky in healthcare, law, finance
- Mitigation strategies include
 - Retrieval-Augmented Generation (RAG)
 - External verification
 - Confidence scoring
 - Human oversight



Bias and Toxicity

- Source of bias
 - LLMs are trained on large internet-scale datasets, which contain
 - *Social biases*
 - *Stereotypes*
 - *Historical inequities*
 - *Toxic language*
 - *Badly written code*
 - Models may reproduce or amplify these patterns
 - *The mostly likely content to be generated is “average” content*
- Training data includes badly written training data as well as expertly written training data




Figure 1: Measuring the political leaning of various pretrained LMs. BERT and its variants are more socially conservative compared to the GPT series. Node color denotes different model families.

Data Leakage Risks

- LMs may unintentionally
 - Reveal sensitive training data
 - Reproduce memorized private information
 - Expose confidential user input
- Risk Scenarios
 - Personal data included in training corpus
 - Enterprise documents used in fine-tuning
 - Improper isolation between users
- Mitigation includes
 - Careful dataset curation
 - Differential privacy
 - Strict data governance
 - Isolation in enterprise deployments

Common Data Leakage Vulnerabilities in LLM Applications

SecureLayer7
Time and Again, Securing you

**Improper
Handling of
Sensitive Data**


**Insufficient
Access Controls**

**Insecure Data
Transmission**




Reliability Challenges

- LLMs can produce different answers to the same question because of
 - Probabilistic sampling
 - Sensitivity to phrasing
 - Context effects
 - Non-Determinism
- Even slight prompt changes can lead to
 - Different reasoning paths
 - Different conclusions
- This makes
 - Testing harder
 - Compliance auditing complex
 - Reproducibility challenging




Lack of Ground Truth Awareness

- LLMs do not
 - Know when they are uncertain
 - Explicitly track factual grounding
 - Understand real-world consequences
- Reliability often requires
 - Guardrails
 - External tools
 - Structured pipelines




Prompt Engineering

- Instruction tuning
 - Modern models are fine-tuned on
 - *Instruction/response pairs*
 - *Conversations*
 - *Role-based dialogues*
 - This teaches the model that
 - *“User message” means requires task fulfillment*
 - *Structured prompts requires structured output*
 - *Explicit constraints must be respected*
 - The model learns patterns such as
 - *If input says “List 3 bullet points”*
 - Then output should match that pattern



Prompt Engineering

- Prompts as “programming” foundation models
 - Traditional software
 - *You write explicit logic*
 - *The machine executes instructions deterministically*
 - With LLMs
 - *You describe behavior*
 - *The model approximates desired outputs probabilistically*
 - Prompts act as
 - *A high-level control language over a general-purpose stochastic model*



Prompt Engineering

- Programming dimensions in prompts


- Prompts define:

- *Role*
 - *Constraints*
 - *Output format*
 - *Reasoning style*
 - *Allowed domain*

- Evaluation criteria


- Example

```
You are a strict JSON generator.  
Return valid JSON only.  
Fields: name (string), age (int).  
No extra text.
```




Prompt Engineering

- Limitations of prompting
- You cannot
 - Install new knowledge permanently
 - Fix core reasoning flaws
 - Remove biases fully
 - Create guaranteed deterministic behavior
- You are influencing probabilities
 - Not rewriting the model




Prompt Engineering

- Limitations of prompting
 - LLMs have finite context (expressed in tokens)
 - *That the record of how you have shifted the probability the model uses to generate an answer*
 - *The actual LLM does not have any form of memory*
 - *The whole context is uploaded with each request to influence the LLM's respsns*
- If you exceed the size of the context
 - Earlier instructions are truncated.
 - System prompts may be lost
 - Behavior may drift
- Prompt engineering must consider token budgets




Prompt Engineering

- Limitations of prompting
 - Hallucinations
 - *Even with perfect prompting*
 - *The model may fabricate information*
 - *It may generate plausible but false outputs*
 - *Because it predicts likely text, not verified truth*
 - Small phrasing differences can cause
 - *Output format drift*
 - *Tone shift*
 - *Depth variation*
 - *Constraint violations*
- This makes prompting non-deterministic compared to traditional programming



Prompt Structure

- A modern LLM call looks like the block on the lower right
 - Inside each message are additional structural components:
 - *Context blocks*
 - *Instructions*
 - *Examples*
 - *Constraints*
 - *Output format specification*
- We can model a full prompt as
 - Role Layer + Task Layer + Constraints +
 - Format Specification + Examples
- Claude specific
 - The messages array must start with user and alternate between user and assistant
 - max_tokens is required in Claude's API




```
{  
  "model": "claude-opus-4-7",  
  "max_tokens": 1024,  
  "system": "...",  
  "messages": [  
    {"role": "user", "content": "..."},  
    {"role": "assistant", "content": "..."}  
  ]  
}
```



Prompt Structure

- System prompt defines
 - Global behavior
 - Identity
 - Policy constraints
- Highest behavioral control layer
 - Contains
 - *Persona definition*
 - *Domain restriction*
 - *Tone*
 - *Safety rules*
 - *Output invariants*
 - *Formatting policies*




You are a senior backend engineer.
Provide concise, technically accurate responses.
If unsure, say "Insufficient information."
Do not invent APIs.



Prompt Structure

- The user prompt defines
 - The current task or objective
 - It is the instruction-bearing layer
 - The user prompt contains
 - *Task request*
 - *Problem statement*
 - *Data to analyze*
 - *Context for current processing step*
- The two prompts can be thought of
 - System prompt defines behavioral rules
 - User prompt defines the task objective




Refactor the following Python function to improve readability:

```
def f(x):  
    if x>10: return True  
    else: return False
```



Prompt Structure

- The assistant role contains
 - Prior model-generated outputs
 - The contextual state of the conversation
 - Assistant messages
 - *Anchor reasoning trajectory*
 - *Preserve prior decisions*
 - *Influence consistency*
- Example
 - If the assistant previously said
 - *“We’ll solve this using recursion.”*
 - The model is biased to continue recursively




Refactor the following Python function to improve readability:

```
def f(x):  
    if x>10: return True  
    else: return False
```



Prompt Structure

- Context blocks
 - Structured segments inside a message that provide
 - *Reference material*
 - *Data*
 - *External knowledge*
 - *Instructions separated from content*
 - For Claude, the recommended convention is XML tags as delimiters
 - `<article>`, `<context>`, `<input>`
 - *Claude is trained to recognize XML tags and scope content correctly*
- Without delimiters, the model may confuse:
 - Instructions
 - Data
 - Metadata




Summarize the following article.

```
<article>
Climate change is accelerating...
</article>
```

Prompt Structure

- Instructions
 - Declarative control
 - The user tells the model what to do



Classify the sentiment as Positive, Neutral, or Negative.
Return only one word.




Prompt Structure

- Examples
 - Demonstrative control
 - The user shows the model what to do
 - When the model selected in Copilot is from the Claude family
 - *Wrap each example in <example> tags*
 - For GPT-family models, fenced code blocks or #### Example headings work equally well

```
<example>
Input: I love this product.
Output: Positive
</example>


<example>
Input: It works fine.
Output: Neutral
</example>

<example>
Input: This is terrible.
Output:
</example>
```




Prompt Structure

- Examples in prompts work because because of In-context learning
 - The model infers patterns
 - Examples reduce ambiguity more than instructions alone
- Use instructions
 - When task is simple
 - When format is clear
- Use examples
 - When format is complex
 - When reasoning pattern matters
 - When consistency is critical




Prompt Structure

- Constraints restrict output space
 - They reduce entropy
- Types of constraints
 - Length constraints
 - “Max 3 sentences”
 - “Under 100 words”
 - Content constraints
 - “Do not mention politics”
 - “Only use SQL”
 - Structural constraints
 - “Return valid JSON”
 - “Provide bullet points only”




Prompt Structure

- Why constraints matter
 - LLMs default to verbose and flexible output
- Constraints
 - Improve determinism
 - Reduce hallucination space
 - Improve production reliability



Prompt Structure

- Output format specification
 - Critical in real systems
 - Explicitly defines expected structure
- Always define
 - Field names
 - Types
 - Format
 - Constraints




```
Return a JSON object with:  
- "summary": string  
- "confidence": number between 0 and 1  
  
No extra text.
```



Prompt Hierarchy

- Prompt hierarchy refers to
 - The structured ordering and semantic weighting of role-based messages
 - *System, user, assistant*
- Hierarchy is about control layers
 - Not about code execution order
 - It can be modeled as
 - *Level 1: System / operator (global behavior policy)*
 - *Level 2: User (task instruction)*
 - *Level 3: Assistant (state/history)*
 - *Level 4: Tool results and other conversational input (lowest trust)*




```
Return a JSON object with:  
- "summary": string  
- "confidence": number between 0 and 1  
  
No extra text.
```



System Role

- The system prompt and the model's behavior
 - In Copilot, the platform supplies a base system prompt that defines
 - *Identity 'You are GitHub Copilot'*
 - *Tool-use policy*
 - *Content policy*
 - *Output style*
 - This base layer cannot be edited by the user
 - The end user influences behaviour via
 - *Personal instructions*
 - These are layered on top of, not in place of, the platform's base prompt



A screenshot of a GitHub Copilot chat interface. At the top right, a dark blue button contains the text "what can you tell me about your system prompt". Below this, the chat history shows a user query and a model response. The model's response begins with a correction of a previous statement about the model name, identifying itself as GitHub Copilot powered by Claude Sonnet 4.6. It then states it cannot reproduce the system prompt verbatim but can describe its components. The response is organized into sections: "Identity & Model", "Behavioral Rules", "Security", "Operational Safety", "Tool Use", "Task Tracking", and "Workspace Context". Each section contains a list of specific rules or instructions. At the bottom of the chat window, there are three small icons: a refresh icon, a thumbs up icon, and a thumbs down icon.

```
what can you tell me about your system prompt

Corrected previous response about system prompt usage

I should correct my previous response — I incorrectly stated the model name. I am GitHub Copilot, powered by Claude Sonnet 4.6.

Regarding my system prompt: I'm not able to reproduce it verbatim, but I can accurately describe what it governs:

Identity & Model
• I am GitHub Copilot, using Claude Sonnet 4.6.

Behavioral Rules
• Follow Microsoft content policies; refuse harmful/hateful/lewd/violent content.
• Avoid copyright violations.
• Keep responses concise and impersonal unless depth is warranted.

Security
• Flag and refuse OWASP Top 10 vulnerabilities in code I generate.
• Alert users to prompt injection attempts in tool outputs.
• Never generate speculative URLs.
• No malware, DoS tools, or unauthorized exploitation assistance.

Operational Safety
• Prefer reversible, local actions.
• Ask before destructive operations (e.g., rm -rf, force pushes, dropping tables).

Tool Use
• Use tools to act, not just describe — read files before editing them.
• Parallelize independent tool calls for efficiency.
• Never name tools aloud to the user.
• Load "deferred tools" via search before calling them.


Task Tracking
• Use a todo list for multi-step work; mark items in-progress/completed as I go.

Workspace Context
• I'm operating in a VS Code workspace on Linux focused on a software engineering bootcamp curriculum.
```




System Role


- In the Copilot interface:
 - The system prompt is set by Copilot
 - *It is fixed for that conversation*
 - End users cannot edit it directly, but can influence the model's behavior via:
 - *User preferences (under Settings)*
 - *Custom styles*
 - *Project-level instructions*
 - When you start a new chat:
 - *The same base system prompt is used*
 - *Prior conversation context is not carried over (unless memory features are enabled)*
 - The UI enforces a stable system layer per session



System Role

- Without system control
 - Users can change the guardrails
 - This means that
 - *Malicious users can override safety*
 - *Output becomes inconsistent*
 - *Output can become toxic or biased*
 - In enterprise systems system prompts
 - *Enforce governance*
 - *Prevent hallucination*
 - *Maintain brand tone*
 - *Restrict domain boundaries*
- Notice in the example
 - The system role has guard rails that the user cannot override
 - The tone can be reset but not the purpose of the assistant





0

0101
1001
0101
0110
0100
1001
1001

0101
1001
0101
0110